

assignment-2

September 2, 2024

```
[ ]: # Name: Arju Srivastava
     # Section : 2CA
     # Roll No: 12
```

```
[3]: # Load the dataset using Pandas.
import pandas as pd
file = "C:\\Users\\windows\\Downloads\\sensor_data.csv"
df = pd.read_csv(file)
# Display the first few rows of the dataset to understand its structure.
print(df.head())
```

	Timestamp	Temperature	Pressure	Humidity	Vibration
0	2024-01-01 00:00:00	27.483571	NaN	NaN	0.585502
1	2024-01-01 01:00:00	24.308678	101114.677339	55.607845	0.719909
2	2024-01-01 02:00:00	28.238443	101153.642742	NaN	1.373647
3	2024-01-01 03:00:00	32.615149	100923.861365	NaN	1.305185
4	2024-01-01 04:00:00	23.829233	NaN	36.223306	10.000000

```
[7]: # Identify and handle missing values:
     # Replace missing values using the mean or median.

df['Temperature'].fillna(df['Temperature'].mean(), inplace=True)
df['Pressure'].fillna(df['Pressure'].mean(), inplace=True)
df['Humidity'].fillna(df['Humidity'].mean(), inplace=True)
df['Vibration'].fillna(df['Vibration'].mean(), inplace=True)
df
```

C:\Users\windows\AppData\Local\Temp\ipykernel_4664\2308645533.py:4:

FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Temperature'].fillna(df['Temperature'].mean(), inplace=True)
```

C:\Users\windows\AppData\Local\Temp\ipykernel_4664\2308645533.py:5:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Pressure'].fillna(df['Pressure'].mean(), inplace=True)
```

C:\Users\windows\AppData\Local\Temp\ipykernel_4664\2308645533.py:6:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Humidity'].fillna(df['Humidity'].mean(), inplace=True)
```

C:\Users\windows\AppData\Local\Temp\ipykernel_4664\2308645533.py:7:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Vibration'].fillna(df['Vibration'].mean(), inplace=True)
```

```
[7]:
```

	Timestamp	Temperature	Pressure	Humidity	Vibration
0	2024-01-01 00:00:00	27.483571	101889.818616	49.961066	0.585502
1	2024-01-01 01:00:00	24.308678	101114.677339	55.607845	0.719909
2	2024-01-01 02:00:00	28.238443	101153.642742	49.961066	1.373647
3	2024-01-01 03:00:00	32.615149	100923.861365	49.961066	1.305185

4	2024-01-01 04:00:00	23.829233	101889.818616	36.223306	10.000000
..
95	2024-01-04 23:00:00	17.682425	101517.658690	43.070904	0.765412
96	2024-01-05 00:00:00	26.480601	100883.071282	58.995999	0.143433
97	2024-01-05 01:00:00	26.305276	101401.862553	53.072995	1.676936
98	2024-01-05 02:00:00	25.334767	101354.104359	58.128621	0.942730
99	2024-01-05 03:00:00	23.827064	100753.514851	56.296288	1.618908

[100 rows x 5 columns]

```
[9]: # Identify and remove or cap outliers in the dataset using NumPy.
```

```
import numpy as np

def cap_outliers(series, lower_quantile=0.05, upper_quantile=0.95):
    lower_bound = series.quantile(lower_quantile)
    upper_bound = series.quantile(upper_quantile)
    return np.clip(series, lower_bound, upper_bound)

df['Temperature'] = cap_outliers(df['Temperature'])
df['Pressure'] = cap_outliers(df['Pressure'])
df['Humidity'] = cap_outliers(df['Humidity'])
df['Vibration'] = cap_outliers(df['Vibration'])
```

```
[11]: # Convert the Timestamp column to a proper datetime format.
```

```
df['Timestamp'] = pd.to_datetime(df['Timestamp'])
df
```

```
[11]:
```

	Timestamp	Temperature	Pressure	Humidity	Vibration
0	2024-01-01 00:00:00	27.483571	101889.818616	49.961066	0.585502
1	2024-01-01 01:00:00	24.308678	101114.677339	55.607845	0.719909
2	2024-01-01 02:00:00	28.238443	101153.642742	49.961066	1.373647
3	2024-01-01 03:00:00	32.390034	100923.861365	49.961066	1.305185
4	2024-01-01 04:00:00	23.829233	101889.818616	36.223306	1.915785
..
95	2024-01-04 23:00:00	17.682425	101517.658690	43.070904	0.765412
96	2024-01-05 00:00:00	26.480601	100883.071282	58.995999	0.496599
97	2024-01-05 01:00:00	26.305276	101401.862553	53.072995	1.676936
98	2024-01-05 02:00:00	25.334767	101354.104359	58.128621	0.942730
99	2024-01-05 03:00:00	23.827064	100753.514851	56.296288	1.618908

[100 rows x 5 columns]

```
[14]: # Normalize the Temperature, Pressure, and Vibration columns to a range between
# 0 and 1 using Min-Max scaling.
```

```
from sklearn.preprocessing import MinMaxScaler
```

```

scaler = MinMaxScaler()

df[['Temperature', 'Pressure', 'Vibration']] = scaler.
    ↳fit_transform(df[['Temperature', 'Pressure', 'Vibration']])
df

```

```

[14]:
      Timestamp  Temperature  Pressure  Humidity  Vibration
0  2024-01-01 00:00:00      0.693635  0.764077  49.961066   0.062644
1  2024-01-01 01:00:00      0.495391  0.280652  55.607845   0.157351
2  2024-01-01 02:00:00      0.740770  0.304953  49.961066   0.617994
3  2024-01-01 03:00:00      1.000000  0.161647  49.961066   0.569753
4  2024-01-01 04:00:00      0.465454  0.764077  36.223306   1.000000
..          ...          ...          ...          ...
95 2024-01-04 23:00:00      0.081641  0.531975  43.070904   0.189414
96 2024-01-05 00:00:00      0.631009  0.136208  58.995999   0.000000
97 2024-01-05 01:00:00      0.620061  0.459758  53.072995   0.831700
98 2024-01-05 02:00:00      0.559462  0.429973  58.128621   0.314357
99 2024-01-05 03:00:00      0.465319  0.055409  56.296288   0.790812

[100 rows x 5 columns]

```

```

[19]: # Create a new column that calculates the moving average of Temperature over a
    ↳window of 10 readings.
df['Temperature_Moving_Avg'] = df['Temperature'].rolling(window=10).mean()
df.tail()

```

```

[19]:
      Timestamp  Temperature  Pressure  Humidity  Vibration \
95 2024-01-04 23:00:00      0.081641  0.531975  43.070904   0.189414
96 2024-01-05 00:00:00      0.631009  0.136208  58.995999   0.000000
97 2024-01-05 01:00:00      0.620061  0.459758  53.072995   0.831700
98 2024-01-05 02:00:00      0.559462  0.429973  58.128621   0.314357
99 2024-01-05 03:00:00      0.465319  0.055409  56.296288   0.790812

      Temperature_Moving_Avg
95                0.534409
96                0.515075
97                0.512962
98                0.531591
99                0.508243

```

```

[16]: # Plot the time series data for Temperature, Pressure, Humidity, and Vibration.
import matplotlib.pyplot as plt

plt.figure(figsize=(14, 8))

plt.subplot(2, 2, 1)
plt.plot(df['Timestamp'], df['Temperature'], label='Temperature')

```

```

plt.xlabel('Timestamp')
plt.ylabel('Temperature (C)')
plt.title('Temperature Time Series')


plt.subplot(2, 2, 2)
plt.plot(df['Timestamp'], df['Pressure'], label='Pressure', color='orange')
plt.xlabel('Timestamp')
plt.ylabel('Pressure (Pa)')
plt.title('Pressure Time Series')

plt.subplot(2, 2, 3)
plt.plot(df['Timestamp'], df['Humidity'], label='Humidity', color='green')
plt.xlabel('Timestamp')
plt.ylabel('Humidity (%)')
plt.title('Humidity Time Series')

plt.subplot(2, 2, 4)
plt.plot(df['Timestamp'], df['Vibration'], label='Vibration', color='red')
plt.xlabel('Timestamp')
plt.ylabel('Vibration (mm/s)')
plt.title('Vibration Time Series')

plt.tight_layout()
plt.show()


```



```
[21]: # Create histograms for each sensor to visualize the distribution of the
      ↪ readings.
      df[['Temperature', 'Pressure', 'Humidity', 'Vibration']].hist(bins=30,
      ↪ figsize=(14,8))

[20]: # Plot a heatmap to visualize the correlation between different sensor readings.
      import seaborn as sns

      plt.figure(figsize=(10, 6))
      correlation_matrix = df[['Temperature', 'Pressure', 'Humidity', 'Vibration']].
      ↪ corr()
      sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
      plt.title("Correlation HeatMap")
      plt.show()
```



```
[ ]:
```